

# WRF Scaling and Performance Assessment

## Comparison of Compilers and MPI Libraries on Cheyenne

---

Akira Kyle<sup>1</sup>, Davide Del Vento<sup>2</sup>, Brian Vanderwende<sup>2</sup>, Negin Sobhani<sup>2</sup>,  
Dixit Patel<sup>3</sup>

August 2, 2018

<sup>1</sup>**Carnegie Mellon University**



- Background
  - WRF
  - Cheyenne
  - Benchmark Cases
- Compilers
- Message Passing Interface Libraries
- Run Time Scaling
- Computation Time Scaling
- MVAPICH scaling

# Background

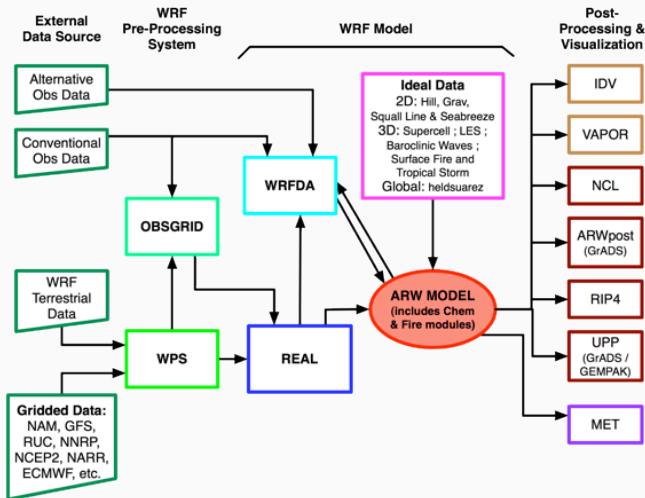
---

- The Weather Research and Forecast (WRF) model is a parallel mesoscale numerical weather forecasting application used in both operational and research environments.

- The Weather Research and Forecast (WRF) model is a parallel mesoscale numerical weather forecasting application used in both operational and research environments.
- WRF is among the more commonly run codes by atmospheric scientists on NCAR's Cheyenne supercomputer.

- The Weather Research and Forecast (WRF) model is a parallel mesoscale numerical weather forecasting application used in both operational and research environments.
- WRF is among the more commonly run codes by atmospheric scientists on NCAR's Cheyenne supercomputer.
  - Thus it is very important for WRF's users to know how to obtain the best performance of WRF on Cheyenne, especially as users scale their runs to larger core counts.

## WRF Modeling System Flow Chart



- 4,032 computation nodes



- 4,032 computation nodes
  - Dual-socket nodes, 18 cores per socket

- 4,032 computation nodes
  - Dual-socket nodes, 18 cores per socket
    - 145,152 total processor cores

- 4,032 computation nodes
  - Dual-socket nodes, 18 cores per socket
    - 145,152 total processor cores
  - 2.3-GHz Intel Xeon E5-2697V4 (Broadwell) processors

- 4,032 computation nodes
  - Dual-socket nodes, 18 cores per socket
    - 145,152 total processor cores
  - 2.3-GHz Intel Xeon E5-2697V4 (Broadwell) processors
    - 16 flops per clock

- 4,032 computation nodes
  - Dual-socket nodes, 18 cores per socket
    - 145,152 total processor cores
  - 2.3-GHz Intel Xeon E5-2697V4 (Broadwell) processors
    - 16 flops per clock
  - 5.34 peak petaflops

- 4,032 computation nodes
  - Dual-socket nodes, 18 cores per socket
    - 145,152 total processor cores
  - 2.3-GHz Intel Xeon E5-2697V4 (Broadwell) processors
    - 16 flops per clock
  - 5.34 peak petaflops
- 313 TB total system memory

- 4,032 computation nodes
  - Dual-socket nodes, 18 cores per socket
    - 145,152 total processor cores
  - 2.3-GHz Intel Xeon E5-2697V4 (Broadwell) processors
    - 16 flops per clock
  - 5.34 peak petaflops
- 313 TB total system memory
  - 64 GB/node on 3,168 nodes, DDR4-2400

- 4,032 computation nodes
  - Dual-socket nodes, 18 cores per socket
    - 145,152 total processor cores
  - 2.3-GHz Intel Xeon E5-2697V4 (Broadwell) processors
    - 16 flops per clock
  - 5.34 peak petaflops
- 313 TB total system memory
  - 64 GB/node on 3,168 nodes, DDR4-2400
  - 128 GB/node on 864 nodes, DDR4-2400



- 4,032 computation nodes
  - Dual-socket nodes, 18 cores per socket
    - 145,152 total processor cores
  - 2.3-GHz Intel Xeon E5-2697V4 (Broadwell) processors
    - 16 flops per clock
  - 5.34 peak petaflops
- 313 TB total system memory
  - 64 GB/node on 3,168 nodes, DDR4-2400
  - 128 GB/node on 864 nodes, DDR4-2400
- Mellanox EDR InfiniBand high-speed interconnect

- 4,032 computation nodes
  - Dual-socket nodes, 18 cores per socket
    - 145,152 total processor cores
  - 2.3-GHz Intel Xeon E5-2697V4 (Broadwell) processors
    - 16 flops per clock
  - 5.34 peak petaflops
- 313 TB total system memory
  - 64 GB/node on 3,168 nodes, DDR4-2400
  - 128 GB/node on 864 nodes, DDR4-2400
- Mellanox EDR InfiniBand high-speed interconnect
  - Partial 9D Enhanced Hypercube single-plane interconnect topology

- 4,032 computation nodes
  - Dual-socket nodes, 18 cores per socket
    - 145,152 total processor cores
  - 2.3-GHz Intel Xeon E5-2697V4 (Broadwell) processors
    - 16 flops per clock
  - 5.34 peak petaflops
- 313 TB total system memory
  - 64 GB/node on 3,168 nodes, DDR4-2400
  - 128 GB/node on 864 nodes, DDR4-2400
- Mellanox EDR InfiniBand high-speed interconnect
  - Partial 9D Enhanced Hypercube single-plane interconnect topology
  - Bandwidth: 25 GBps bidirectional per link

- 4,032 computation nodes
  - Dual-socket nodes, 18 cores per socket
    - 145,152 total processor cores
  - 2.3-GHz Intel Xeon E5-2697V4 (Broadwell) processors
    - 16 flops per clock
  - 5.34 peak petaflops
- 313 TB total system memory
  - 64 GB/node on 3,168 nodes, DDR4-2400
  - 128 GB/node on 864 nodes, DDR4-2400
- Mellanox EDR InfiniBand high-speed interconnect
  - Partial 9D Enhanced Hypercube single-plane interconnect topology
  - Bandwidth: 25 GBps bidirectional per link
  - Latency: MPI ping-pong  $< 1 \mu\text{s}$ ; hardware link 130 ns

- Official CONUS (Contiguous United States) benchmarks for WRF only ran on WRF versions 3.8.1 or prior.

- Official CONUS (Contiguous United States) benchmarks for WRF only ran on WRF versions 3.8.1 or prior.
- Wanted to benchmark most recent version of WRF (4.0), so we had to update the old CONUS benchmarks.

- Official CONUS (Contiguous United States) benchmarks for WRF only ran on WRF versions 3.8.1 or prior.
- Wanted to benchmark most recent version of WRF (4.0), so we had to update the old CONUS benchmarks.
  - Also created several new benchmark cases.

- Official CONUS (Contiguous United States) benchmarks for WRF only ran on WRF versions 3.8.1 or prior.
- Wanted to benchmark most recent version of WRF (4.0), so we had to update the old CONUS benchmarks.
  - Also created several new benchmark cases.
- Benchmark cases cover commonly used physics parameterizations.



- Official CONUS (Contiguous United States) benchmarks for WRF only ran on WRF versions 3.8.1 or prior.
- Wanted to benchmark most recent version of WRF (4.0), so we had to update the old CONUS benchmarks.
  - Also created several new benchmark cases.
- Benchmark cases cover commonly used physics parameterizations.
  - CONUS benchmarks use the CONUS physics suite.

- Official CONUS (Contiguous United States) benchmarks for WRF only ran on WRF versions 3.8.1 or prior.
- Wanted to benchmark most recent version of WRF (4.0), so we had to update the old CONUS benchmarks.
  - Also created several new benchmark cases.
- Benchmark cases cover commonly used physics parameterizations.
  - CONUS benchmarks use the CONUS physics suite.
    - But 2.5 km resolution case disables `cu_physics`.

- Official CONUS (Contiguous United States) benchmarks for WRF only ran on WRF versions 3.8.1 or prior.
- Wanted to benchmark most recent version of WRF (4.0), so we had to update the old CONUS benchmarks.
  - Also created several new benchmark cases.
- Benchmark cases cover commonly used physics parameterizations.
  - CONUS benchmarks use the CONUS physics suite.
    - But 2.5 km resolution case disables `cu_physics`.
  - Hurricane Maria benchmarks use the TROPICAL physics suite

- Official CONUS (Contiguous United States) benchmarks for WRF only ran on WRF versions 3.8.1 or prior.
- Wanted to benchmark most recent version of WRF (4.0), so we had to update the old CONUS benchmarks.
  - Also created several new benchmark cases.
- Benchmark cases cover commonly used physics parameterizations.
  - CONUS benchmarks use the CONUS physics suite.
    - But 2.5 km resolution case disables `cu_physics`.
  - Hurricane Maria benchmarks use the TROPICAL physics suite
    - But `cu_physics` disabled and `sf_sfclay_physics` = 1 for both resolutions.

Region	Resolution	Horizontal Gridpoints	Vertical Gridpoints	Total Gridpoints	Time step	Run Hours
CONUS	12 km	425	300	127,500	72	6
CONUS	2.5 km	1901	1301	2,473,201	15	6
Maria	3 km	1396	1384	1,932,064	9	3
Maria	1 km	3665	2894	10,606,510	3	1

# Compilers

---

- GNU Compiler Collection (GCC) versions 6.3.0, 8.1.0

- GNU Compiler Collection (GCC) versions 6.3.0, 8.1.0
  - WRF compiles with `-O2` by default



- GNU Compiler Collection (GCC) versions 6.3.0, 8.1.0
  - WRF compiles with -O2 by default
  - -O3 : enables all -O2 optimization along with optimizations such as function inlining and loop vectorization

- GNU Compiler Collection (GCC) versions 6.3.0, 8.1.0
  - WRF compiles with `-O2` by default
  - `-O3` : enables all `-O2` optimization along with optimizations such as function inlining and loop vectorization
  - `-Ofast` : enables all `-O3` optimizations along with disregarding strict standards compliance (such is for floating point operations)

- GNU Compiler Collection (GCC) versions 6.3.0, 8.1.0
  - WRF compiles with `-O2` by default
  - `-O3` : enables all `-O2` optimization along with optimizations such as function inlining and loop vectorization
  - `-Ofast` : enables all `-O3` optimizations along with disregarding strict standards compliance (such is for floating point operations)
  - `-mfma` : enables Fused Multiply-Add instruction set

- GNU Compiler Collection (GCC) versions 6.3.0, 8.1.0
  - WRF compiles with `-O2` by default
  - `-O3` : enables all `-O2` optimization along with optimizations such as function inlining and loop vectorization
  - `-Ofast` : enables all `-O3` optimizations along with disregarding strict standards compliance (such is for floating point operations)
  - `-mfma` : enables Fused Multiply-Add instruction set
  - `-march=native` : enables target instruction set to be everything supported by the compiling machine


- GNU Compiler Collection (GCC) versions 6.3.0, 8.1.0
  - WRF compiles with `-O2` by default
  - `-O3` : enables all `-O2` optimization along with optimizations such as function inlining and loop vectorization
  - `-Ofast` : enables all `-O3` optimizations along with disregarding strict standards compliance (such is for floating point operations)
  - `-mfma` : enables Fused Multiply-Add instruction set
  - `-march=native` : enables target instruction set to be everything supported by the compiling machine
- Intel Compiler versions 17.0.1, 18.0.1 (Default compiler for Cheyenne)

- GNU Compiler Collection (GCC) versions 6.3.0, 8.1.0
  - WRF compiles with `-O2` by default
  - `-O3` : enables all `-O2` optimization along with optimizations such as function inlining and loop vectorization
  - `-Ofast` : enables all `-O3` optimizations along with disregarding strict standards compliance (such is for floating point operations)
  - `-mfma` : enables Fused Multiply-Add instruction set
  - `-march=native` : enables target instruction set to be everything supported by the compiling machine
- Intel Compiler versions 17.0.1, 18.0.1 (Default compiler for Cheyenne)
  - WRF compiles with `-O3` by default

- GNU Compiler Collection (GCC) versions 6.3.0, 8.1.0
  - WRF compiles with `-O2` by default
  - `-O3` : enables all `-O2` optimization along with optimizations such as function inlining and loop vectorization
  - `-Ofast` : enables all `-O3` optimizations along with disregarding strict standards compliance (such is for floating point operations)
  - `-mfma` : enables Fused Multiply-Add instruction set
  - `-march=native` : enables target instruction set to be everything supported by the compiling machine
- Intel Compiler versions 17.0.1, 18.0.1 (Default compiler for Cheyenne)
  - WRF compiles with `-O3` by default
  - `-xhost` : similar to GNU's `-march=native`

- GNU Compiler Collection (GCC) versions 6.3.0, 8.1.0
  - WRF compiles with `-O2` by default
  - `-O3` : enables all `-O2` optimization along with optimizations such as function inlining and loop vectorization
  - `-Ofast` : enables all `-O3` optimizations along with disregarding strict standards compliance (such is for floating point operations)
  - `-mfma` : enables Fused Multiply-Add instruction set
  - `-march=native` : enables target instruction set to be everything supported by the compiling machine
- Intel Compiler versions 17.0.1, 18.0.1 (Default compiler for Cheyenne)
  - WRF compiles with `-O3` by default
  - `-xhost` : similar to GNU's `-march=native`
  - `-fp-model fast=2` : similar to GNU's `-Ofast` optimization





**Fig. 1:** Comparison of Intel 18.0.1 and Gnu 8.1.0 compilers with various compilation flags normalized to default Intel WRF compilation  
Runs made using CONUS 12 km Benchmark Case on 2 Nodes

- Intel compiler is consistently 25-30% faster than the Gnu compiler across all flags tried.
- We also see that for both Intel and Gnu, the `-Ofast` (for Gnu) or `-fp-model fast=2` (for Intel) are the only flags that make a significant difference in speed.
- Other flags tried such as `-mfma` or `-march=native -Xhost` made little to no difference in WRF's speed.

# **Message Passing Interface Libraries**

---

- SGI's MPT version 2.18 (Default MPI for Cheyenne)


- SGI's MPT version 2.18 (Default MPI for Cheyenne)
- Ohio State University's MVAPICH version 2.2

- SGI's MPT version 2.18 (Default MPI for Cheyenne)
- Ohio State University's MVAPICH version 2.2
- OpenMPI version 3.1.0


- SGI's MPT version 2.18 (Default MPI for Cheyenne)
- Ohio State University's MVAPICH version 2.2
- OpenMPI version 3.1.0
- Intel MPI version 2018.1.163

- SGI's MPT version 2.18 (Default MPI for Cheyenne)
- Ohio State University's MVAPICH version 2.2
- OpenMPI version 3.1.0
- Intel MPI version 2018.1.163
- MPICH version 3.2





**Fig. 2:** MPI comparison using  
**Gnu 8.1.0**




**Fig. 3:** MPI comparison using  
**Intel 18.0.1**

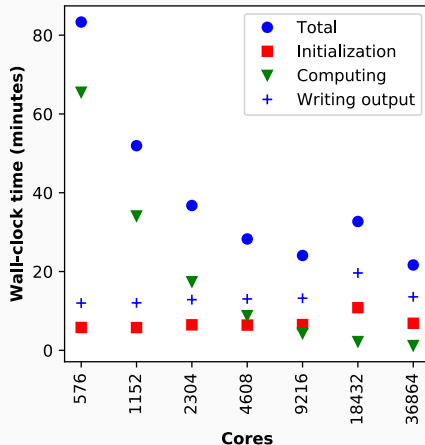
- MPT, MVAPICH and OpenMPI all have similar performance.
- MPICH has overall poor performance and the performance.
- Intel MPI does not scale well to large node counts.

# Total Run Time Scaling

---



**Fig. 4:** Run Time Scaling on Yellowstone



**Fig. 5:** Run Time Scaling on Cheyenne

Runs made using Hurricane Maria 1 km Benchmark case.

- In Fig 5 for Cheyenne the initialization and writing output times remain relatively fixed, only increasing slightly as you move to larger core counts.


- In Fig 5 for Cheyenne the initialization and writing output times remain relatively fixed, only increasing slightly as you move to larger core counts.
- In Fig 4 for Yellowstone, the initialization time scaled much poorer at large node counts, eventually leading to unfeasible long jobs.

- In Fig 5 for Cheyenne the initialization and writing output times remain relatively fixed, only increasing slightly as you move to larger core counts.
- In Fig 4 for Yellowstone, the initialization time scaled much poorer at large node counts, eventually leading to unfeasible long jobs.
  - This improvement in the scaling of the initialization time is likely due to the improvements made to WRF's initialization code with how the MPI calls are performed along with improvements in the MPI used in Cheyenne versus Yellowstone.

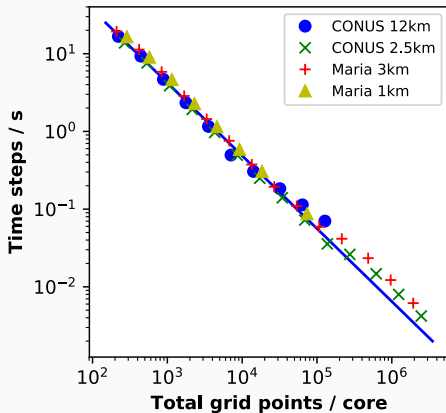
# Computation Time Scaling

---





**Fig. 6:** Computation Scaling on  
Yellowstone



**Fig. 7:** Computation Scaling on  
Cheyenne

- Large number of gridpoints per core region:

- Large number of gridpoints per core region:
  - In both Fig 6 on Yellowstone and Fig 7 on Cheyenne, WRF experiences linear **strong scaling**

- Large number of gridpoints per core region:
  - In both Fig 6 on Yellowstone and Fig 7 on Cheyenne, WRF experiences linear **strong scaling**
    - Increasing number of cores will proportionately decrease computation time while the same number of total core-hours will be used for computation

- Large number of gridpoints per core region:
  - In both Fig 6 on Yellowstone and Fig 7 on Cheyenne, WRF experiences linear **strong scaling**
    - Increasing number of cores will proportionately decrease computation time while the same number of total core-hours will be used for computation
- Small number of gridpoints per core region:

- Large number of gridpoints per core region:
  - In both Fig 6 on Yellowstone and Fig 7 on Cheyenne, WRF experiences linear **strong scaling**
    - Increasing number of cores will proportionately decrease computation time while the same number of total core-hours will be used for computation
- Small number of gridpoints per core region:
  - In Fig 6 on Yellowstone, WRF departs from the linear strong scaling relationship.

- Large number of gridpoints per core region:
  - In both Fig 6 on Yellowstone and Fig 7 on Cheyenne, WRF experiences linear **strong scaling**
    - Increasing number of cores will proportionately decrease computation time while the same number of total core-hours will be used for computation
- Small number of gridpoints per core region:
  - In Fig 6 on Yellowstone, WRF departs from the linear strong scaling relationship.
    - Runs in this region would use more core-hours to run the same simulation than if they had been run on fewer cores

- Large number of gridpoints per core region:
  - In both Fig 6 on Yellowstone and Fig 7 on Cheyenne, WRF experiences linear **strong scaling**
    - Increasing number of cores will proportionately decrease computation time while the same number of total core-hours will be used for computation
- Small number of gridpoints per core region:
  - In Fig 6 on Yellowstone, WRF departs from the linear strong scaling relationship.
    - Runs in this region would use more core-hours to run the same simulation than if they had been run on fewer cores
    - MPI communication dominates the actual time spent in computation



- Large number of gridpoints per core region:
  - In both Fig 6 on Yellowstone and Fig 7 on Cheyenne, WRF experiences linear **strong scaling**
    - Increasing number of cores will proportionately decrease computation time while the same number of total core-hours will be used for computation
- Small number of gridpoints per core region:
  - In Fig 6 on Yellowstone, WRF departs from the linear strong scaling relationship.
    - Runs in this region would use more core-hours to run the same simulation than if they had been run on fewer cores
    - MPI communication dominates the actual time spent in computation
  - In Fig 7 on Cheyenne, WRF does not significantly depart from linear strong scaling

- Large number of gridpoints per core region:
  - In both Fig 6 on Yellowstone and Fig 7 on Cheyenne, WRF experiences linear **strong scaling**
    - Increasing number of cores will proportionately decrease computation time while the same number of total core-hours will be used for computation
- Small number of gridpoints per core region:
  - In Fig 6 on Yellowstone, WRF departs from the linear strong scaling relationship.
    - Runs in this region would use more core-hours to run the same simulation than if they had been run on fewer cores
    - MPI communication dominates the actual time spent in computation
  - In Fig 7 on Cheyenne, WRF does not significantly depart from linear strong scaling
    - Likely due to improvements in WRF's MPI code along and a better network interconnect on Cheyenne than Yellowstone

- Starting with V4.0, WRF refuses to run with a minimum patch size of less than 10 grid points in either direction

- Starting with V4.0, WRF refuses to run with a minimum patch size of less than 10 grid points in either direction
  - Prevents users from running with fewer than 100 gridpoints per core where WRF computation would be very MPI bound

- Starting with V4.0, WRF refuses to run with a minimum patch size of less than 10 grid points in either direction
  - Prevents users from running with fewer than 100 gridpoints per core where WRF computation would be very MPI bound
- Normal nodes on Cheyenne have only 64 GB of memory, less than on Yellowstone

- Starting with V4.0, WRF refuses to run with a minimum patch size of less than 10 grid points in either direction
  - Prevents users from running with fewer than 100 gridpoints per core where WRF computation would be very MPI bound
- Normal nodes on Cheyenne have only 64 GB of memory, less than on Yellowstone
  - WRF runs with too many gridpoints per node will run out of memory and be killed

- Starting with V4.0, WRF refuses to run with a minimum patch size of less than 10 grid points in either direction
  - Prevents users from running with fewer than 100 gridpoints per core where WRF computation would be very MPI bound
- Normal nodes on Cheyenne have only 64 GB of memory, less than on Yellowstone
  - WRF runs with too many gridpoints per node will run out of memory and be killed
    - Typically the max gridpoints per node that will fit into memory is between  $10^5$  and  $10^6$  total gridpoints but it depends on the physics parameterizations

- Starting with V4.0, WRF refuses to run with a minimum patch size of less than 10 grid points in either direction
  - Prevents users from running with fewer than 100 gridpoints per core where WRF computation would be very MPI bound
- Normal nodes on Cheyenne have only 64 GB of memory, less than on Yellowstone
  - WRF runs with too many gridpoints per node will run out of memory and be killed
    - Typically the max gridpoints per node that will fit into memory is between  $10^5$  and  $10^6$  total gridpoints but it depends on the physics parameterizations
  - The points in very large gridpoints per core region in Fig 7 for Cheyenne, used the 128 GB memory nodes and undersubscribed the cores on each node.



- Starting with V4.0, WRF refuses to run with a minimum patch size of less than 10 grid points in either direction
  - Prevents users from running with fewer than 100 gridpoints per core where WRF computation would be very MPI bound
- Normal nodes on Cheyenne have only 64 GB of memory, less than on Yellowstone
  - WRF runs with too many gridpoints per node will run out of memory and be killed
    - Typically the max gridpoints per node that will fit into memory is between  $10^5$  and  $10^6$  total gridpoints but it depends on the physics parameterizations
  - The points in very large gridpoints per core region in Fig 7 for Cheyenne, used the 128 GB memory nodes and undersubscribed the cores on each node.
    - Undersubscription of cores is likely responsible for the small bump in speed observed

- Starting with V4.0, WRF refuses to run with a minimum patch size of less than 10 grid points in either direction
  - Prevents users from running with fewer than 100 gridpoints per core where WRF computation would be very MPI bound
- Normal nodes on Cheyenne have only 64 GB of memory, less than on Yellowstone
  - WRF runs with too many gridpoints per node will run out of memory and be killed
    - Typically the max gridpoints per node that will fit into memory is between  $10^5$  and  $10^6$  total gridpoints but it depends on the physics parameterizations
  - The points in very large gridpoints per core region in Fig 7 for Cheyenne, used the 128 GB memory nodes and undersubscribed the cores on each node.
    - Undersubscription of cores is likely responsible for the small bump in speed observed
    - Undersubscription of cores is an inefficient use of a user's core-hour allocation and users should in general not do this

# MVAPICH Scaling

---

- Interested in MVAPICH as a potential future default MPI

- Interested in MVAPICH as a potential future default MPI
- MVAPICH developed for InfiniBand networks

- Interested in MVAPICH as a potential future default MPI
- MVAPICH developed for InfiniBand networks
- Tried setting some runtime environment variables:

- Interested in MVAPICH as a potential future default MPI
- MVAPICH developed for InfiniBand networks
- Tried setting some runtime environment variables:
  - BIND

- Interested in MVAPICH as a potential future default MPI
- MVAPICH developed for InfiniBand networks
- Tried setting some runtime environment variables:
  - BIND
    - `MV2_CPU_BINDING_POLICY=hybrid`




- Interested in MVAPICH as a potential future default MPI
- MVAPICH developed for InfiniBand networks
- Tried setting some runtime environment variables:
  - BIND
    - MV2\_CPU\_BINDING\_POLICY=hybrid
    - MV2\_HYBRID\_BINDING\_POLICY=bunch


- Interested in MVAPICH as a potential future default MPI
- MVAPICH developed for InfiniBand networks
- Tried setting some runtime environment variables:
  - BIND
    - MV2\_CPU\_BINDING\_POLICY=hybrid
    - MV2\_HYBRID\_BINDING\_POLICY=bunch
  - HW

- Interested in MVAPICH as a potential future default MPI
- MVAPICH developed for InfiniBand networks
- Tried setting some runtime environment variables:
  - BIND
    - MV2\_CPU\_BINDING\_POLICY=hybrid
    - MV2\_HYBRID\_BINDING\_POLICY=bunch
  - HW
    - MV2\_USE\_MCAST=1

- Interested in MVAPICH as a potential future default MPI
- MVAPICH developed for InfiniBand networks
- Tried setting some runtime environment variables:
  - BIND
    - MV2\_CPU\_BINDING\_POLICY=hybrid
    - MV2\_HYBRID\_BINDING\_POLICY=bunch
  - HW
    - MV2\_USE\_MCAST=1
    - MV2\_ENABLE\_SHARP=1



**Fig. 8:** MVAPICH CONUS 12 km Init and Write Scaling



**Fig. 9:** MVAPICH Maria 1km Init and Write Scaling

## Conclusion

---

- Intel compiler consistently faster than Gnu compiler
  - Users should use `-fp-model fast=2` or `-Ofast` for a modest performance increase
- MPT, OpenMPI, and MVAPICH show similar performance while Intel MPI and MPICH have poorer performance
- WRF's initialization and writing time show improvements compared to previous results on Yellowstone with a previous WRF version
- WRF V4.0 scales well across entire run-able region
  - Will run out of memory on runs with too many of gridpoints per core
  - WRF will prevent runs with too few of gridpoints per core



- Davide Del Vento
- Brian Vanderwende
- Alessandro Fanfarillo
- Negin Sobhani
- The SIParCS Program
  - AJ Lauer
  - Jenna Preston
  - Elliott Foust
  - Rich Loft
  - Shilo Hall